

Video Agent

One prompt becomes a continuous 40-second story — four 10-second shots with enforced narrative and visual continuity.

The problem

Text-to-video models generate clips of 5–10 seconds in isolation. Generate four clips from four prompts and you get four unrelated clips: the protagonist changes face, the room changes colour, the story never moves. Generation is solved; **continuity is not**.

How it works

- 1

Plan the story — one LLM pass produces a 4-beat arc (setup, development, turn, resolution) summing to exactly 40s.
- 2

Lock a continuity bible — canonical character, wardrobe, location, lighting, palette and lens language. Immutable for the life of the job.
- 3

Generate shots sequentially — via Higgsfield MCP behind a provider abstraction. Each prompt = bible + beat action + camera move.
- 4

Chain the frames — the final frame of shot n conditions shot n+1, so identity carries forward.
- 5

QC and repair — a vision model scores each shot against the bible; failures regenerate that shot only, capped at 2 attempts.
- 6

Assemble and deliver — ffmpeg stitch, normalise, optional music bed, presigned URLs.

Deliberate trade-off

Shots run **sequentially, not in parallel**. Parallel is roughly 4× faster but breaks frame chaining, and frame chaining is what makes the product work. Latency was traded for the core value proposition.

Delivery

M1–M2	Job lifecycle, planning, continuity bible
M3	Higgsfield MCP, frame chaining, assembly
M4	QC loop, partial results, resume
M5	Observability, cost caps, load + chaos

Success metrics

METRIC	V1 TARGET
Story coherence (human, 1–5)	≥ 4.0
Jobs with continuity score ≥ 0.75	≥ 85%
p90 end-to-end job latency	≤ 8 min
Jobs failing with zero deliverable	< 1%

Resilience

- Never returns nothing — if one shot succeeded, a stitched partial is delivered with a working [resume](#)
- Resume, don't restart — completed shots are never regenerated or re-billed
- Shot-level regeneration — fix shot 3, leave 1, 2 and 4 byte-identical
- Provider abstraction — capability negotiation plus failover, so an API change is not an outage

Key risks

RISK	MITIGATION
Provider can't hold identity across clips	Frame chaining + locked bible + QC loop
QC itself unreliable → wasted spend	Calibrate on labelled set; cap attempts
Repair loops blow the budget	Hard USD cap; no-progress detection

What's delivered

- The stitched 40-second MP4, plus each 10-second clip separately
- Thumbnail and the extracted continuity frames
- StoryPlan and ContinuityBible as machine-readable JSON
- Per-shot cost, model, seed and prompt — every job is reproducible

Out of scope (v1)

Dialogue and lip-sync · durations other than 40s · user-supplied reference characters · voiceover · editing timeline · above 1080p.